Back-Office Web Traffic on the Internet

Enric Pujol	TU-Berlin
Philipp Richter	TU-Berlin
Balakrishnan Chandrasekaran	Duke University
Georgios Smaragdakis	MIT / TU-Berlin / Akamai
Anja Feldmann	TU-Berlin
Bruce Maggs	Duke University / Akamai
Keung-Chi Ng	Akamai

IMC 2014

Vancouver, BC, CANADA

November 5-7, 2014

The Web for an end user

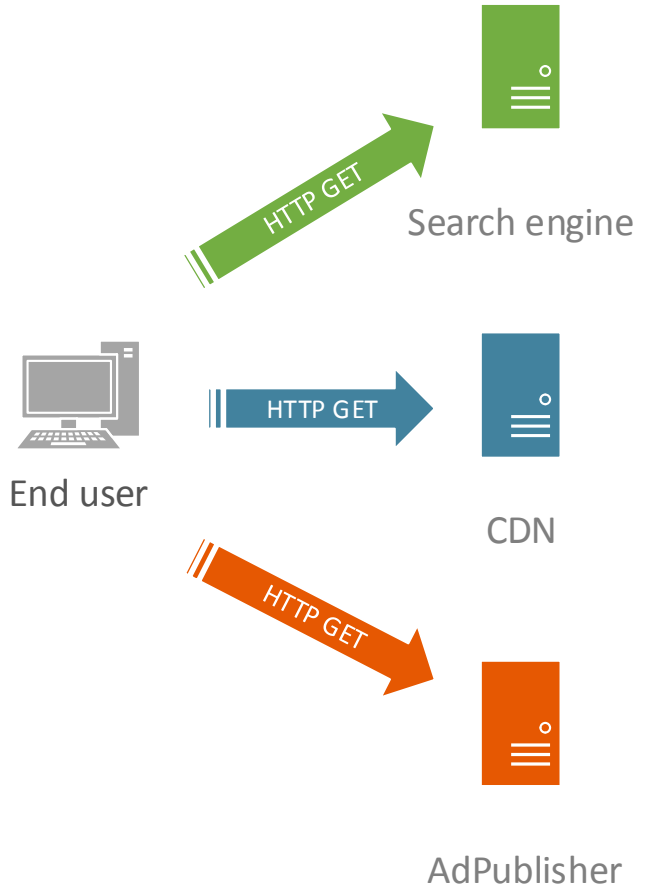

Front-office Web traffic:
Web traffic between end users and servers

Behind the scenes...

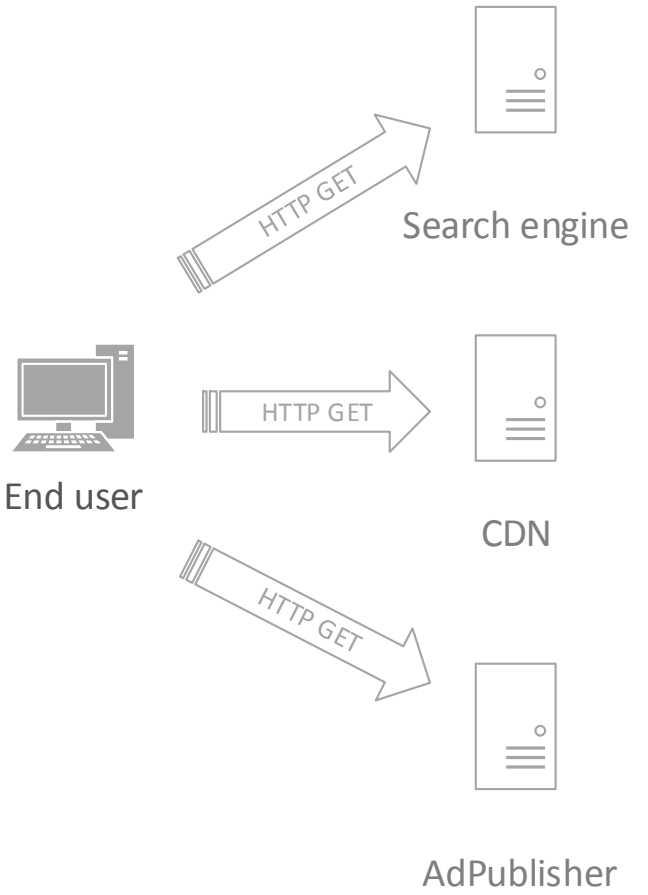

Back-office Web traffic:
Machine-to-machine Web traffic

?

The front-office

The back-office

Search engines: crawlers

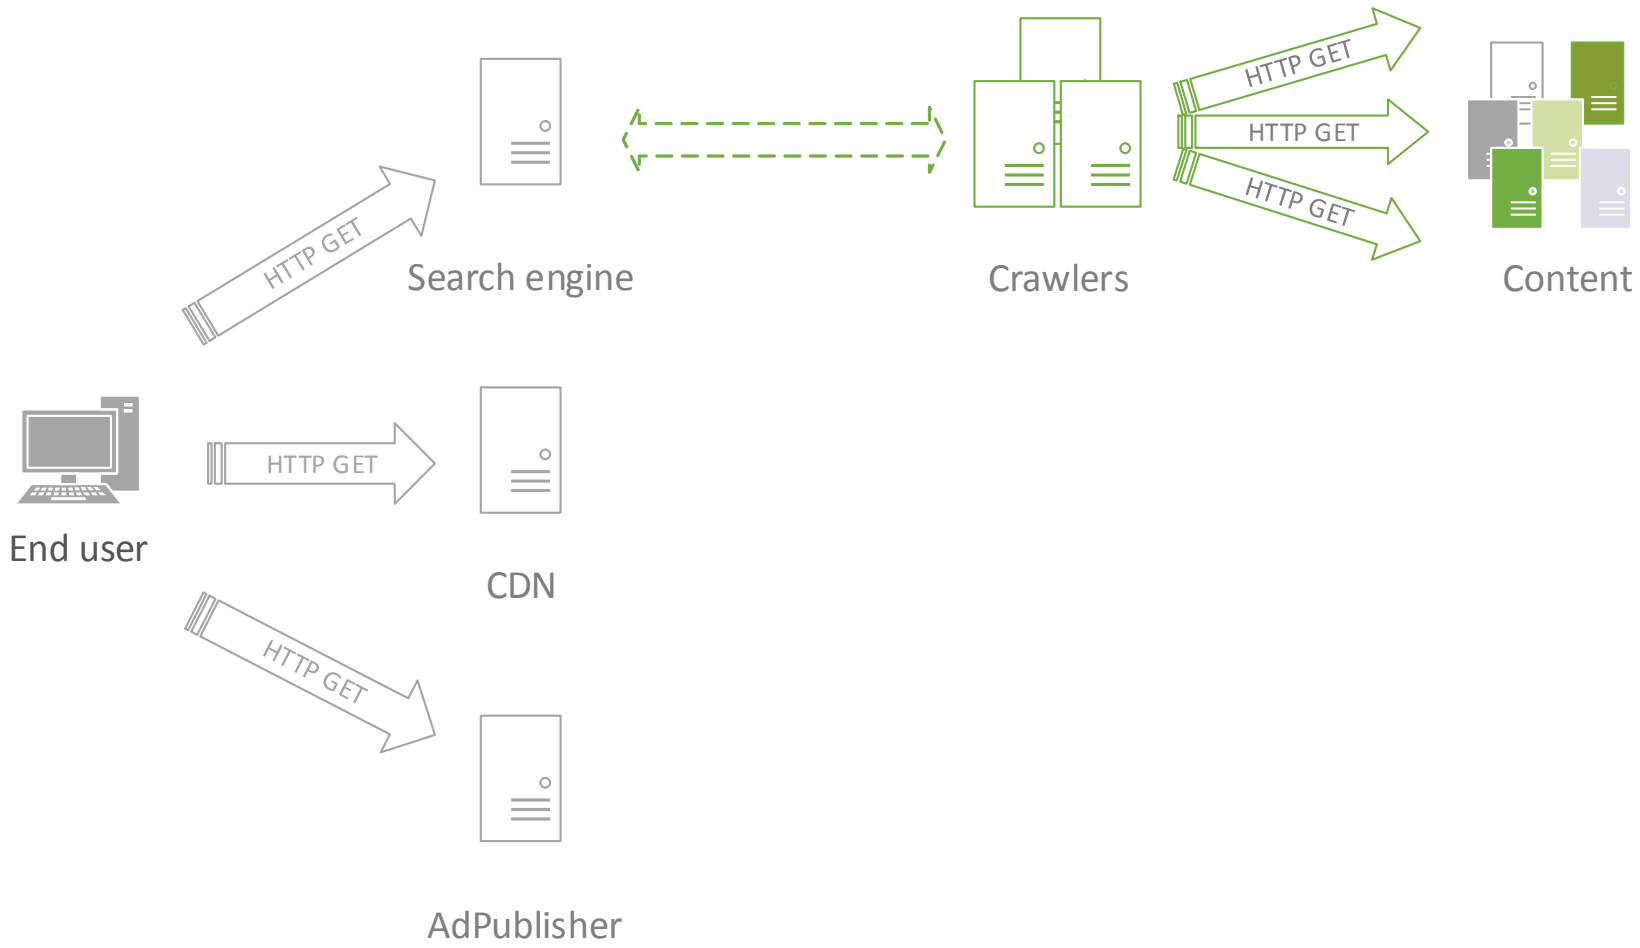

Content delivery: proxies

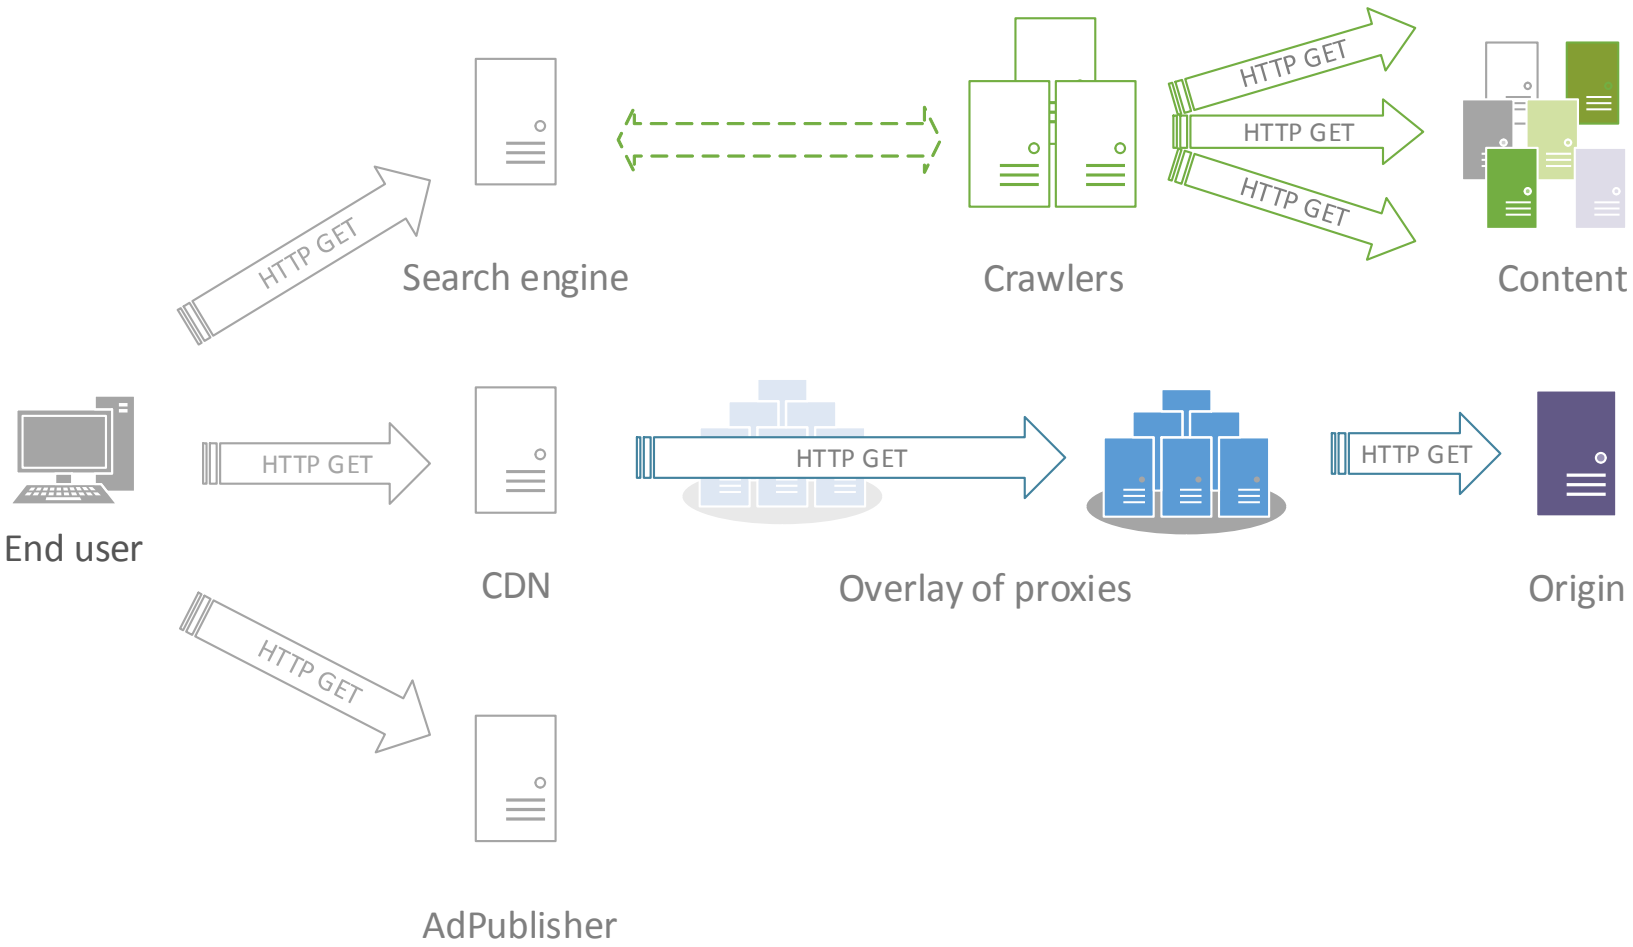

AdExchanges: real-time bidding

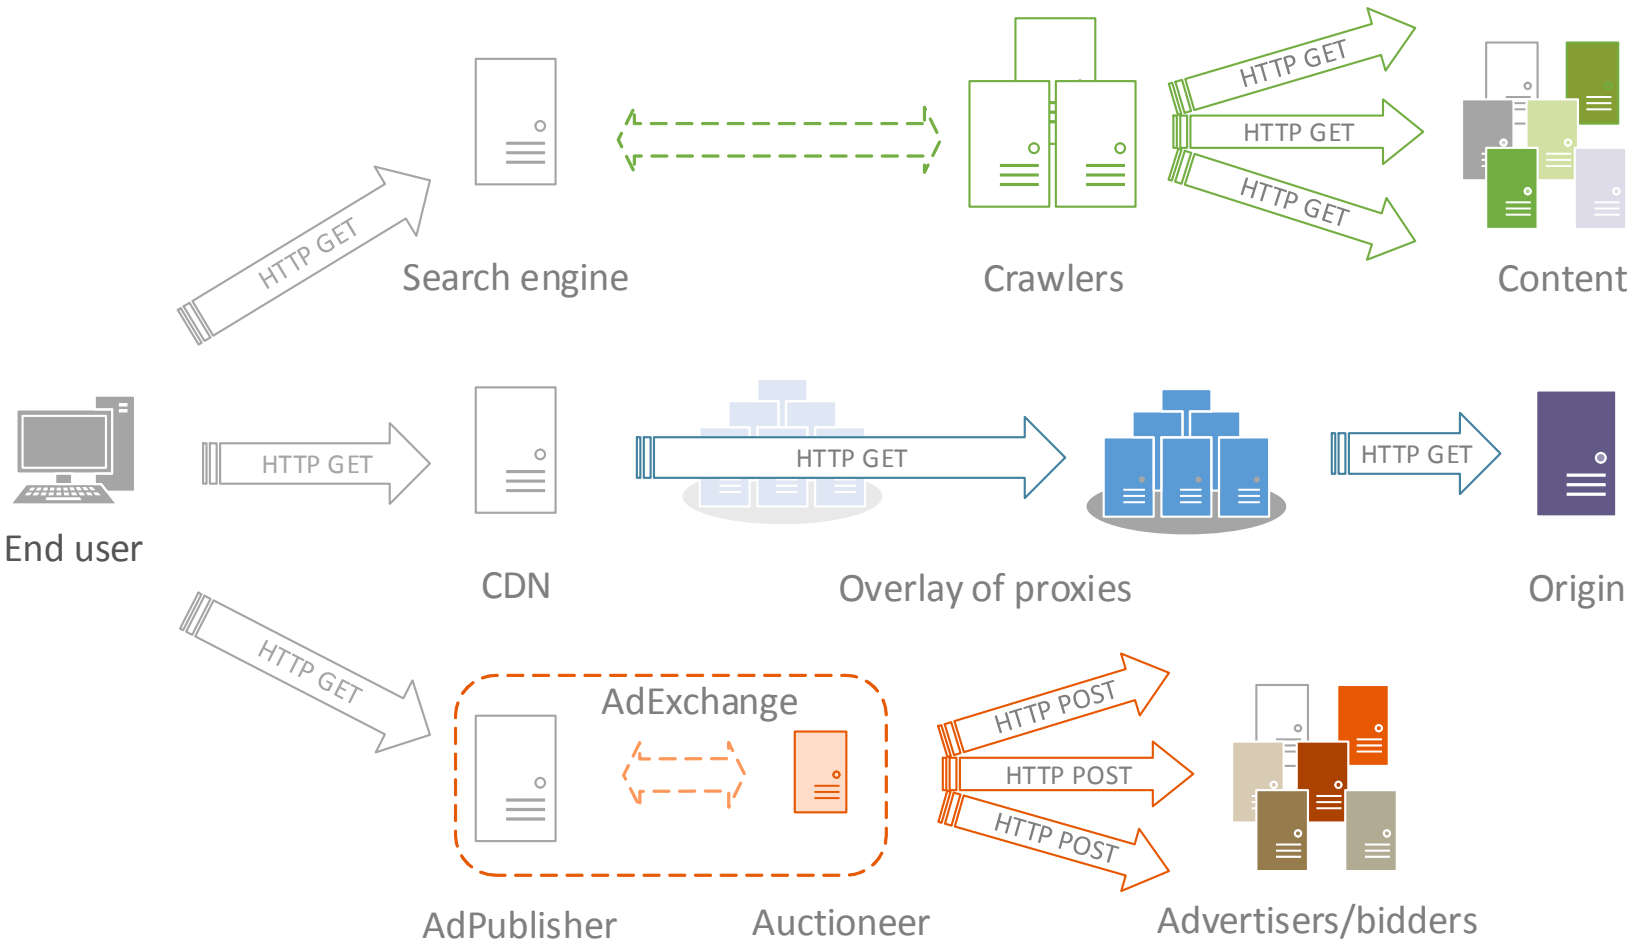

Agenda

1. Introduction
2. Methodology and datasets
3. Characteristics
    1. Traffic
    2. Patterns
    3. Inter-domain perspective
4. CDN back-office traffic
5. The end-user perspective
6. Summary and implications

Vantage points (VP)

Type	VP	Daily traffic	Observations
IXPs	L-IXP	11,900 TB	SFlow (1/16K)
	M-IXP	1,580 TB	
Transit	BBone-1	40 TB	Packet sampled (1/1K)
	BBone-2	70 TB	
Content	CDN	350 TB	5 locations
Eyeballs	RBN	35 TB	Packet dumps

Diverse vantage points: multiple perspectives

Candidate IPs for the back-office

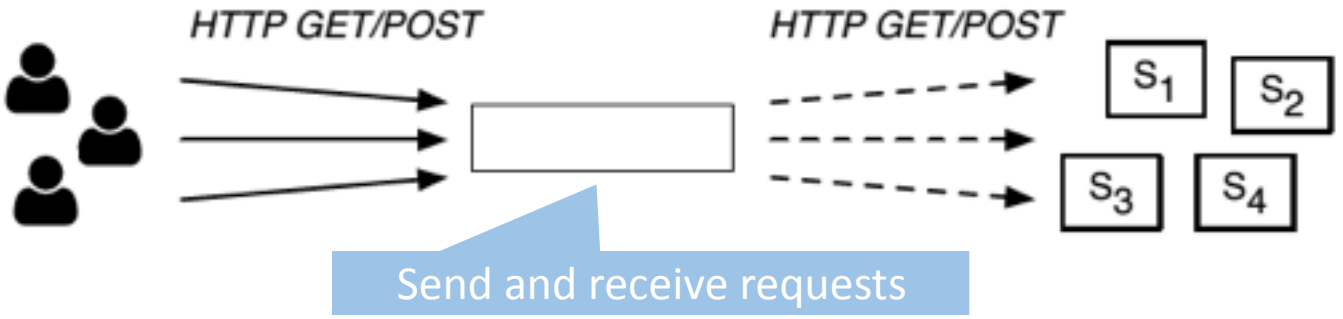

Dual role IPs are prime candidates

Candidate IPs for the back-office

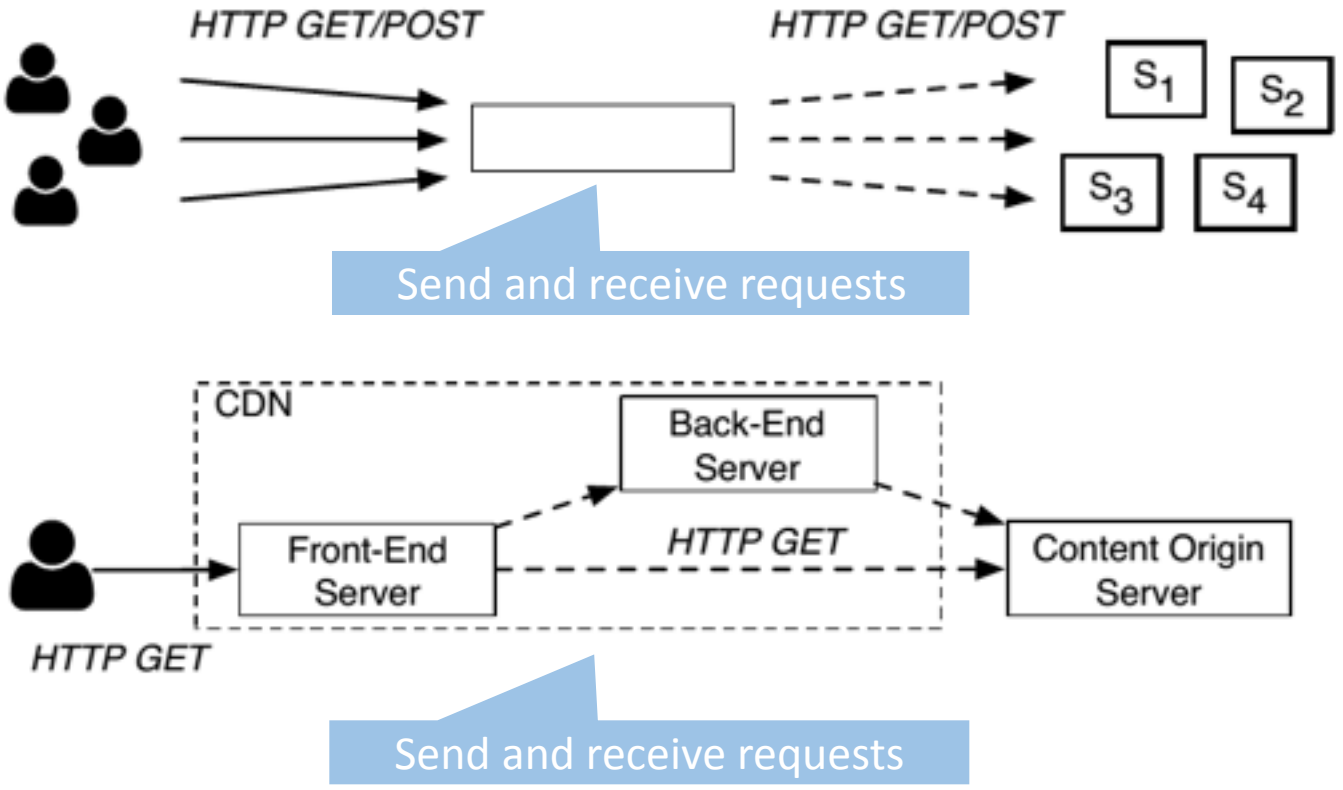

Dual role IPs are prime candidates

Candidate IPs for the back-office

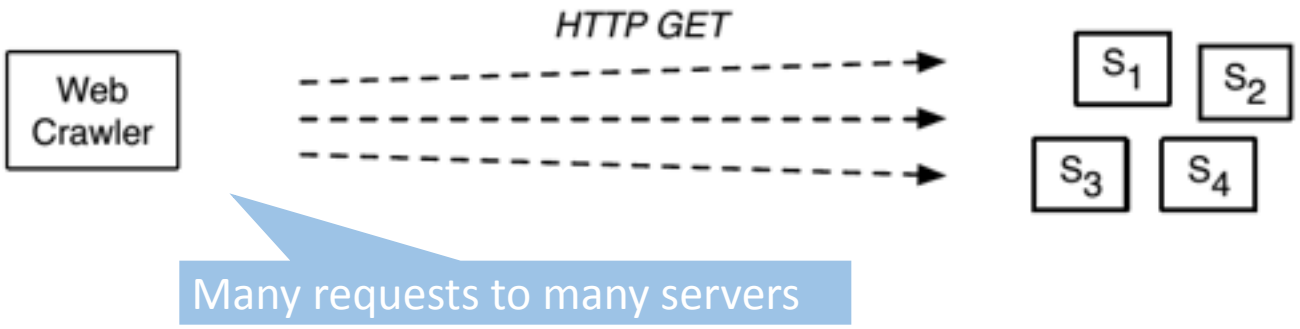

Heavy hitter IPs are also prime candidates

Candidate IPs for the back-office

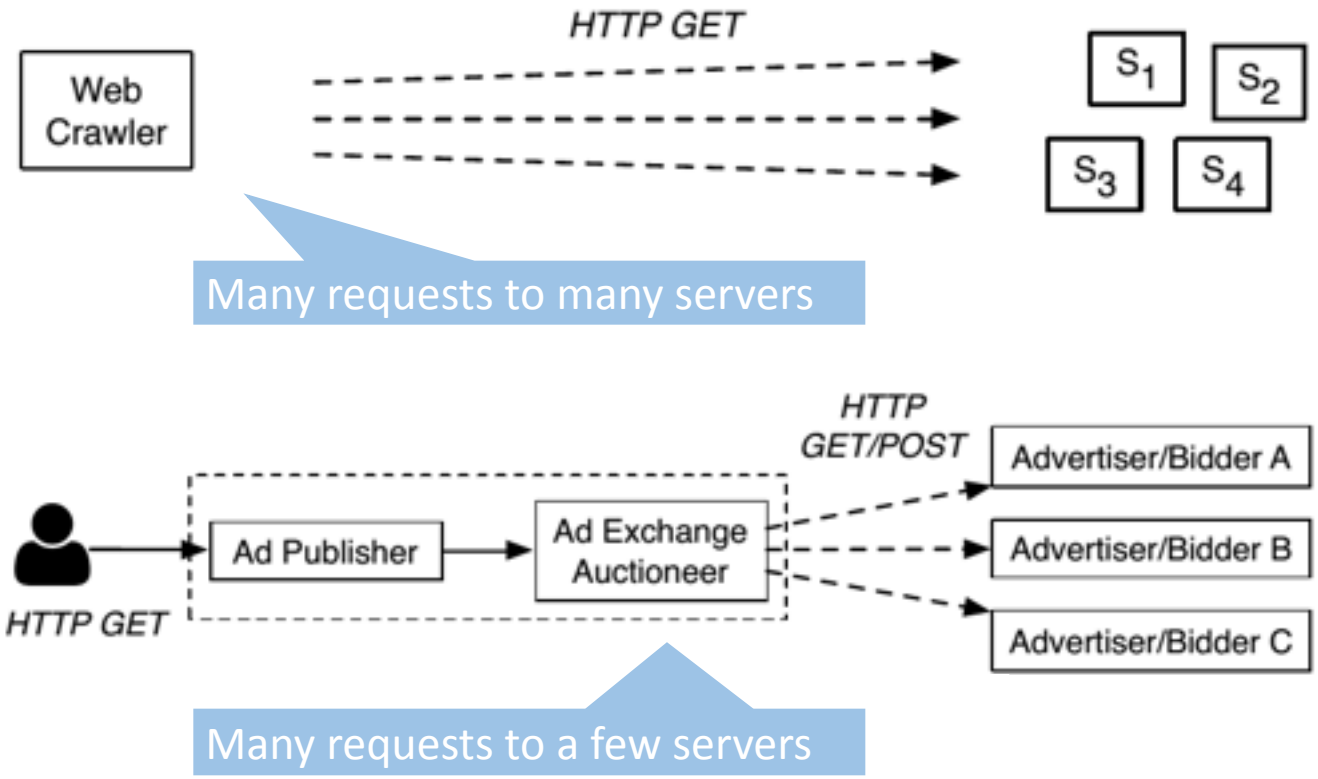

Heavy hitter IPs are also prime candidates

Sources of back-office Web traffic

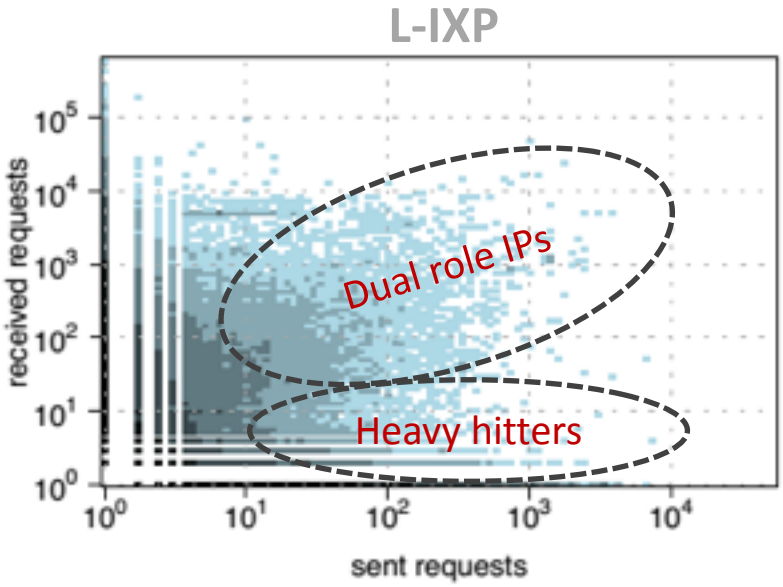

Sources of back-office Web traffic

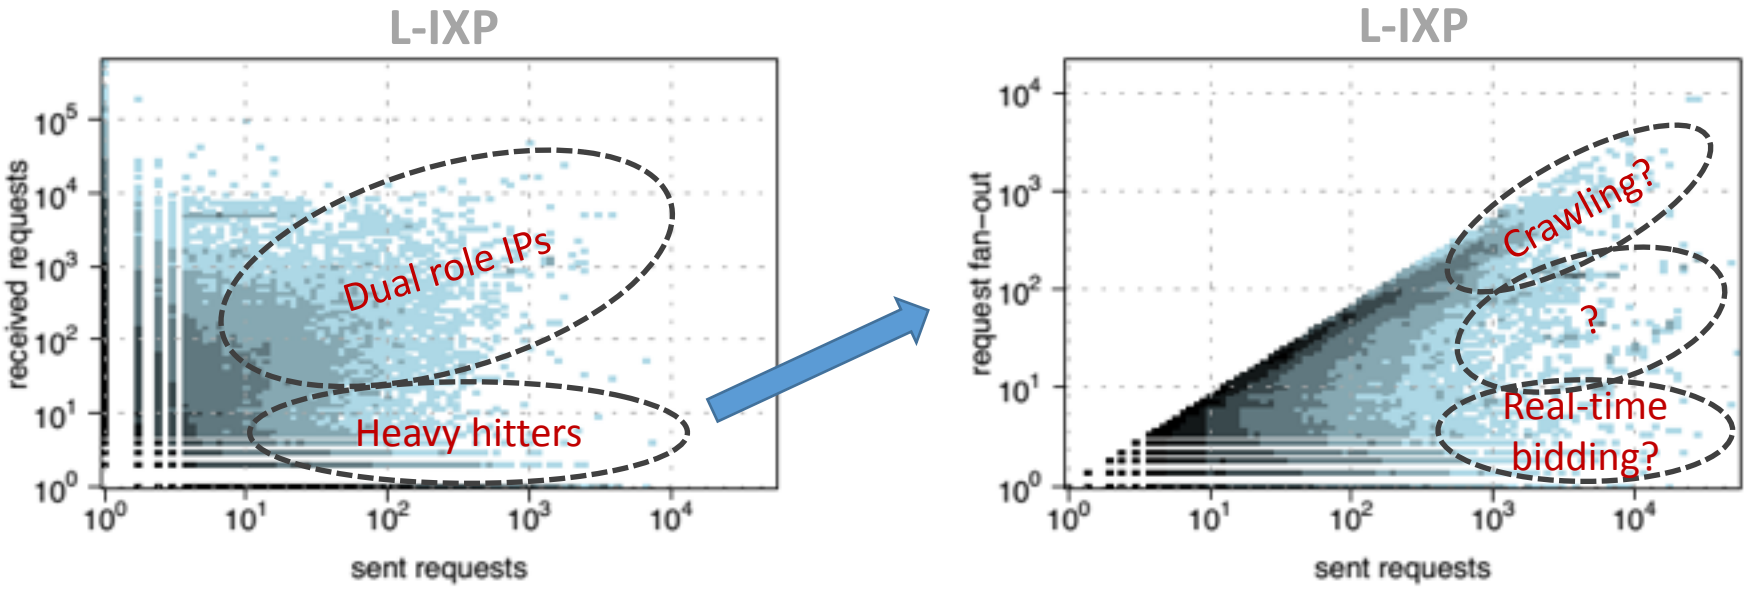

Dual-role IPs: active measurements

		Client only (%)	Server only (%)	Dual-role (%)
L-IXP	**Passive**	96.90	2.74	0.36
	Passive+Active	93.85	2.74	**3.40**

ZMap project: Internet-wide scan of Web Servers (scans.io)

Observations:

1. Most IPs have only client behavior
2. Many servers also show client behavior

Active measurements augment the number of servers

Candidates: manual classification

Crawlers:

- Reverse DNS + Origin AS

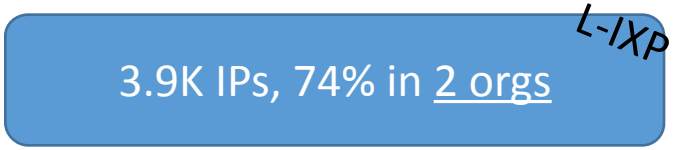

Candidates: manual classification

Crawlers:

- Reverse DNS + Origin AS

3.9K IPs, 74% in 2 orgs

L-IXP

Auctioneers:

- URL + Origin AS

316 IPs, 4 orgs

L-IXP

Candidates: manual classification

Crawlers:

- Reverse DNS + Origin AS

3.9K IPs, 74% in 2 orgs (L-IXP)

Auctioneers:

- URL + Origin AS

316 IPs, 4 orgs (L-IXP)

Content Delivery Proxies:

- Origin AS + Reverse DNS (for caches)

36K IPs, 8 orgs (L-IXP)

Candidates: manual classification

Crawlers:

- Reverse DNS + Origin AS

3.9K IPs, 74% in 2 orgs L-IXP

Auctioneers:

- URL + Origin AS

316 IPs, 4 orgs L-IXP

Content Delivery Proxies:

- Origin AS + Reverse DNS (for caches)

36K IPs, 8 orgs L-IXP

Other:

- Rest of dual-role IPs

151K IPs, mostly in cloud prov. L-IXP

Agenda

1. Introduction
2. Methodology and datasets
3. Characteristics
    1. Traffic
    2. Patterns
    3. Inter-domain perspective
4. CDN back-office traffic
5. The end-user perspective
6. Summary and implications

Traffic

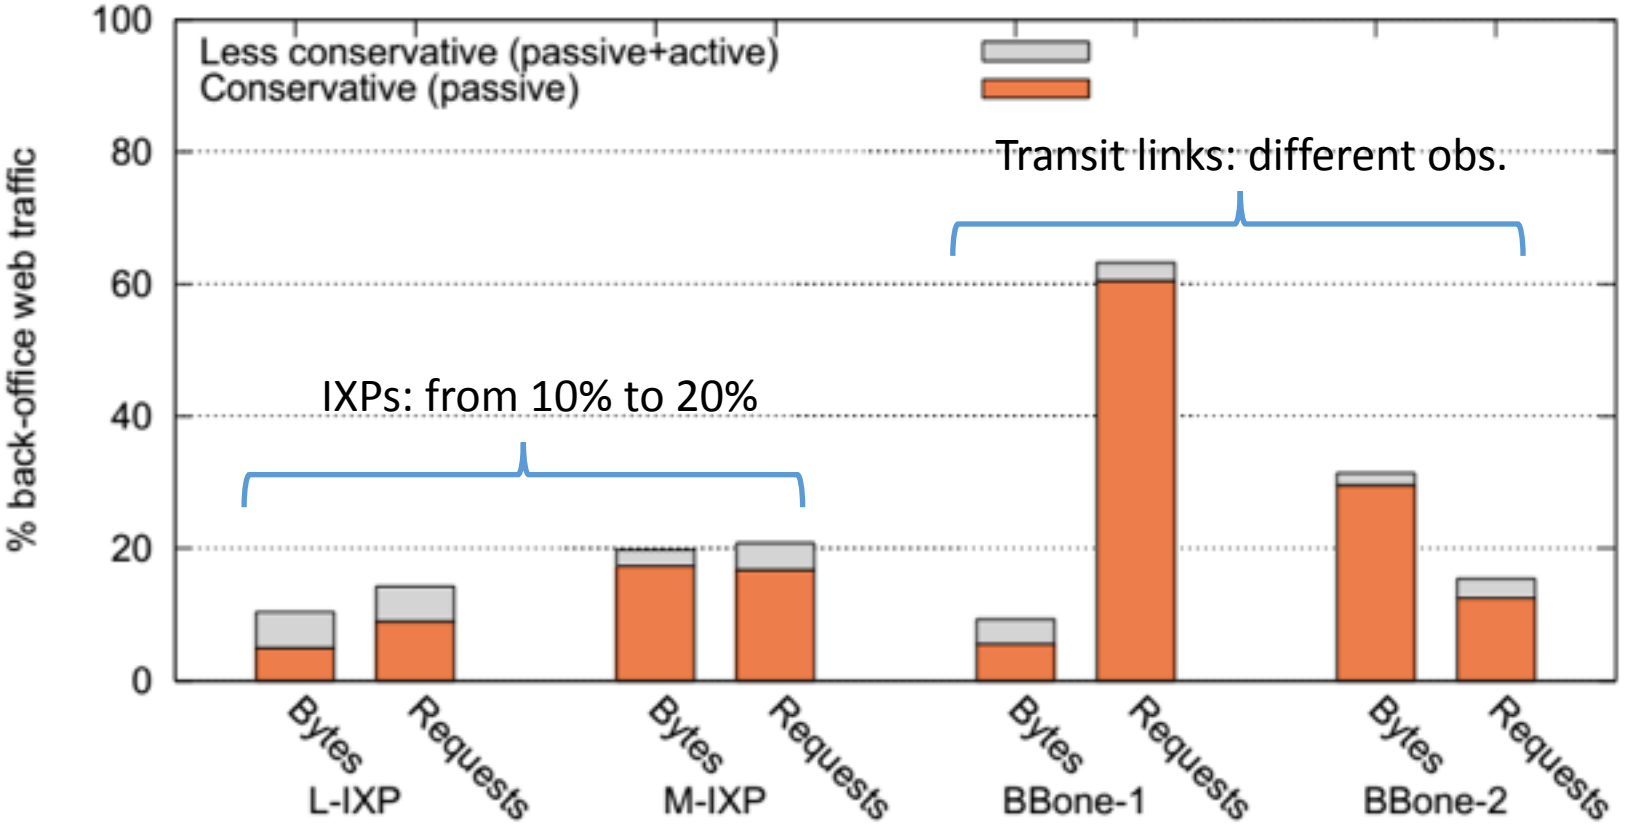

At least 10% in our VPs

Traffic: Contribution per class

		CDPs	Auctioneers	Crawlers	Other
L-IXP	**Bytes**	12.1 %	1.1 %	10.3 %	76.5 %
	Requests	11.8 %	22.5 %	15.1 %	50.6 %

Observations:

1. CDPs — big players – significant share
2. Real-time bidding — many but small transactions
3. Crawlers — a few orgs – significant share
4. Other — cloud service providers

All classes contribute. More to discover

Traffic patterns: bytes

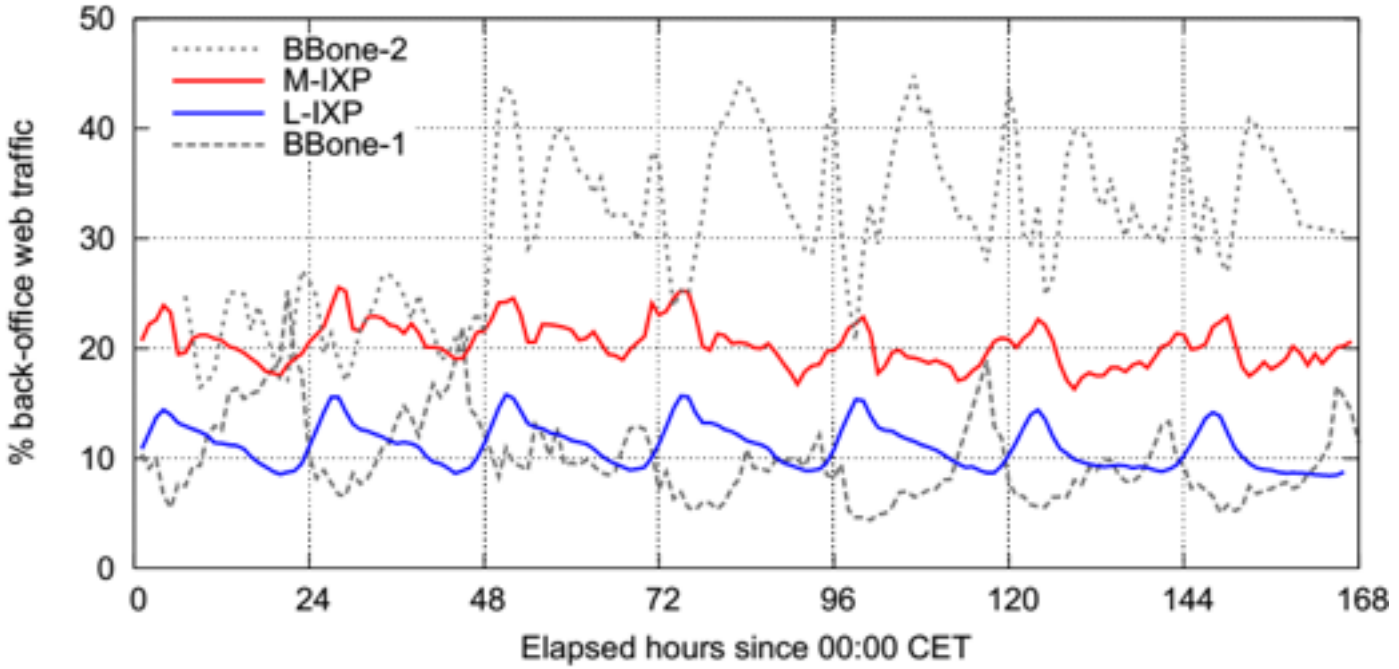

% back-office Web traffic increases during off hours in IXPs

Traffic patterns: requests

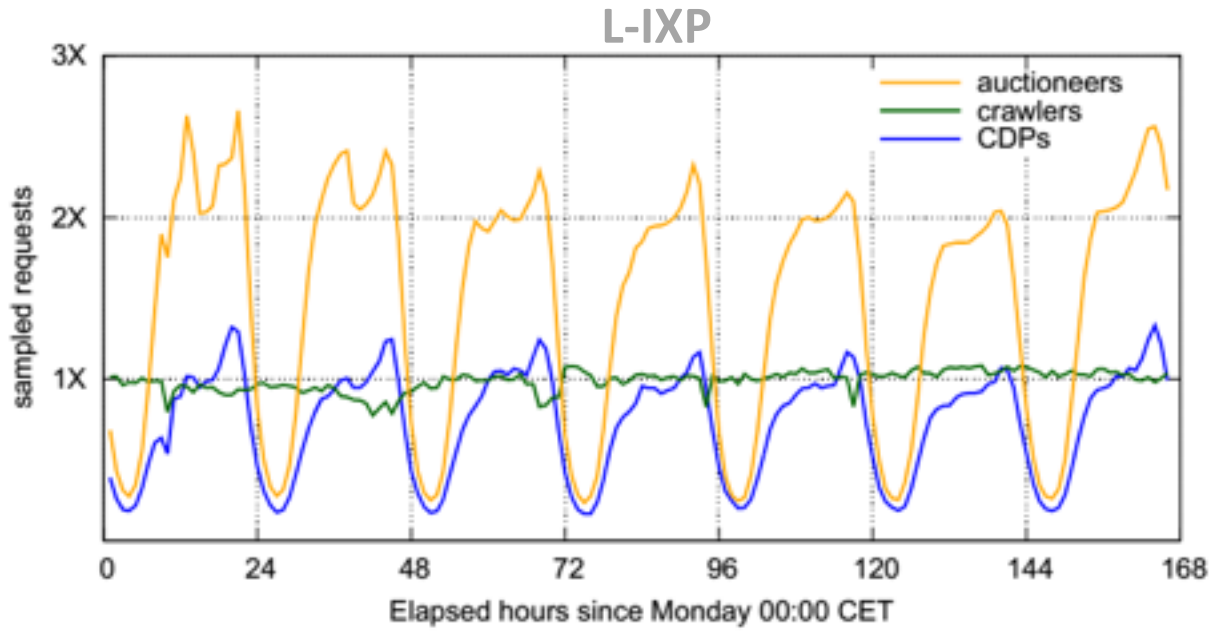

Observations:

1. A multiplicative factor of human activity (e.g., RTB)
2. Non-human triggered activity (e.g., crawlers)

Inter-domain perspective

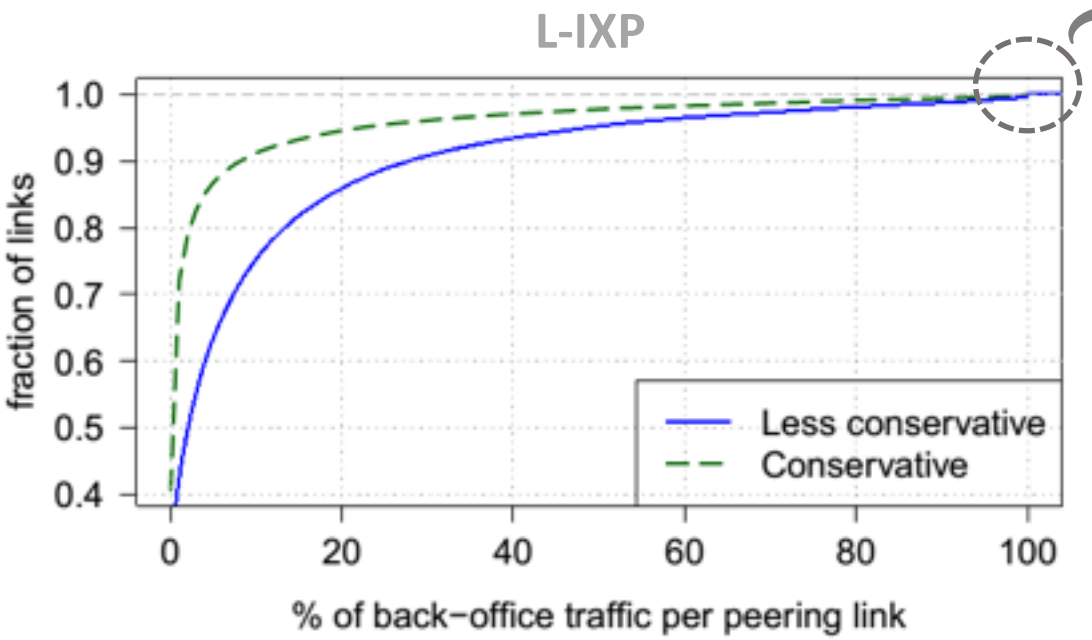

Top 10 traffic carrying links:

4 x Cloud – Content

3 x Search – Hosters

2 x CDN – Content

1 x Content – Advertisement

Back-office traffic appears in many peering links

Agenda

1. Introduction
2. Methodology and datasets
3. Characteristics
    1. Traffic
    2. Patterns
    3. Inter-domain perspective
4. CDN back-office traffic
5. The end-user perspective
6. Summary and implications

A CDN perspective

End-user CDN front-end

Three sub-classes of back-office traffic

A CDN perspective

End-user CDN front-end CDN back-end

Public: front-end back-end over the Internet

A CDN perspective

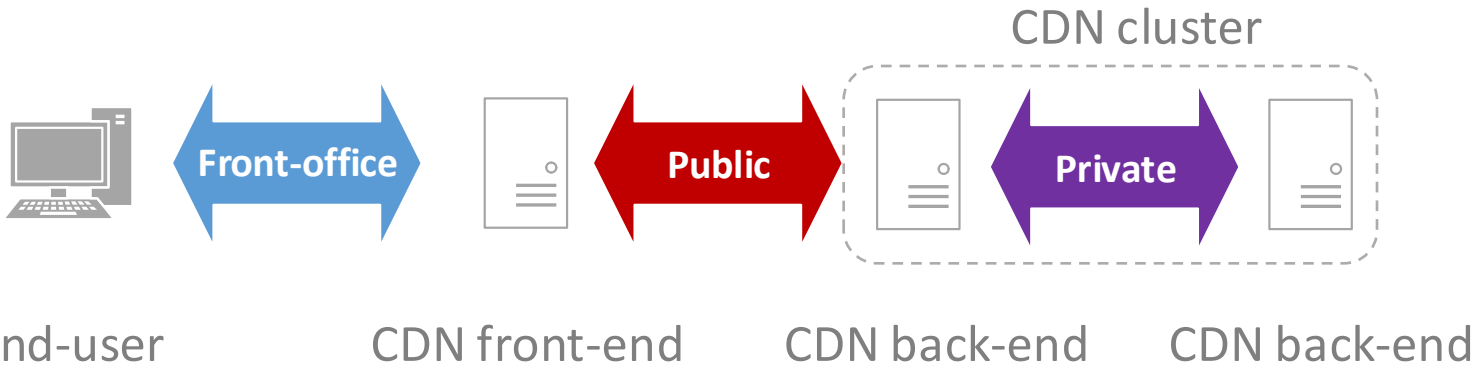

Private: within same cluster

A CDN perspective

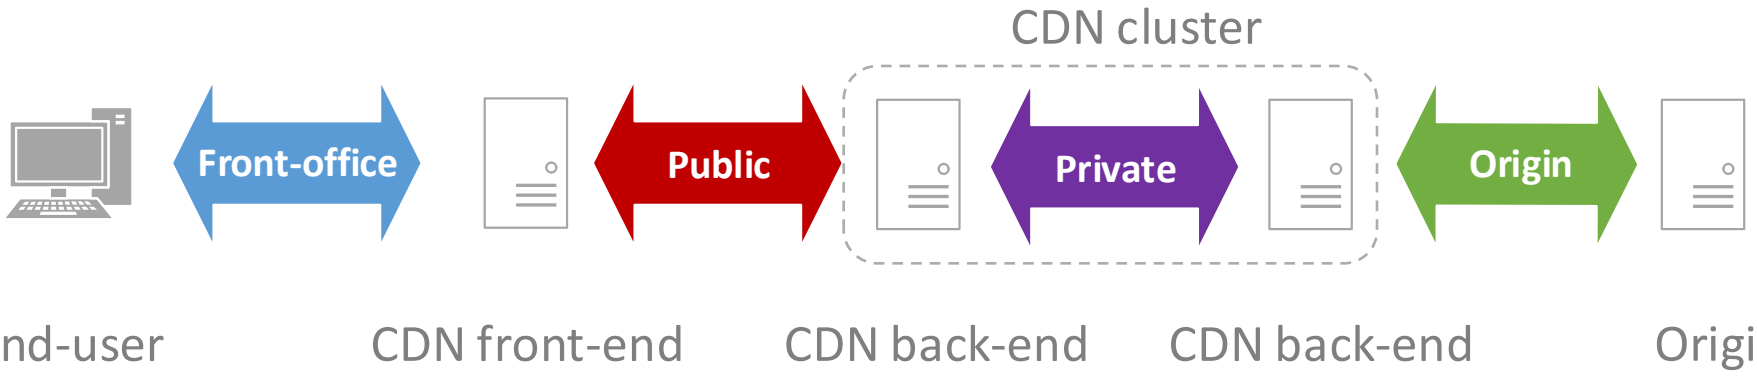

Origin: inter-organization over the Internet

Back-office per location

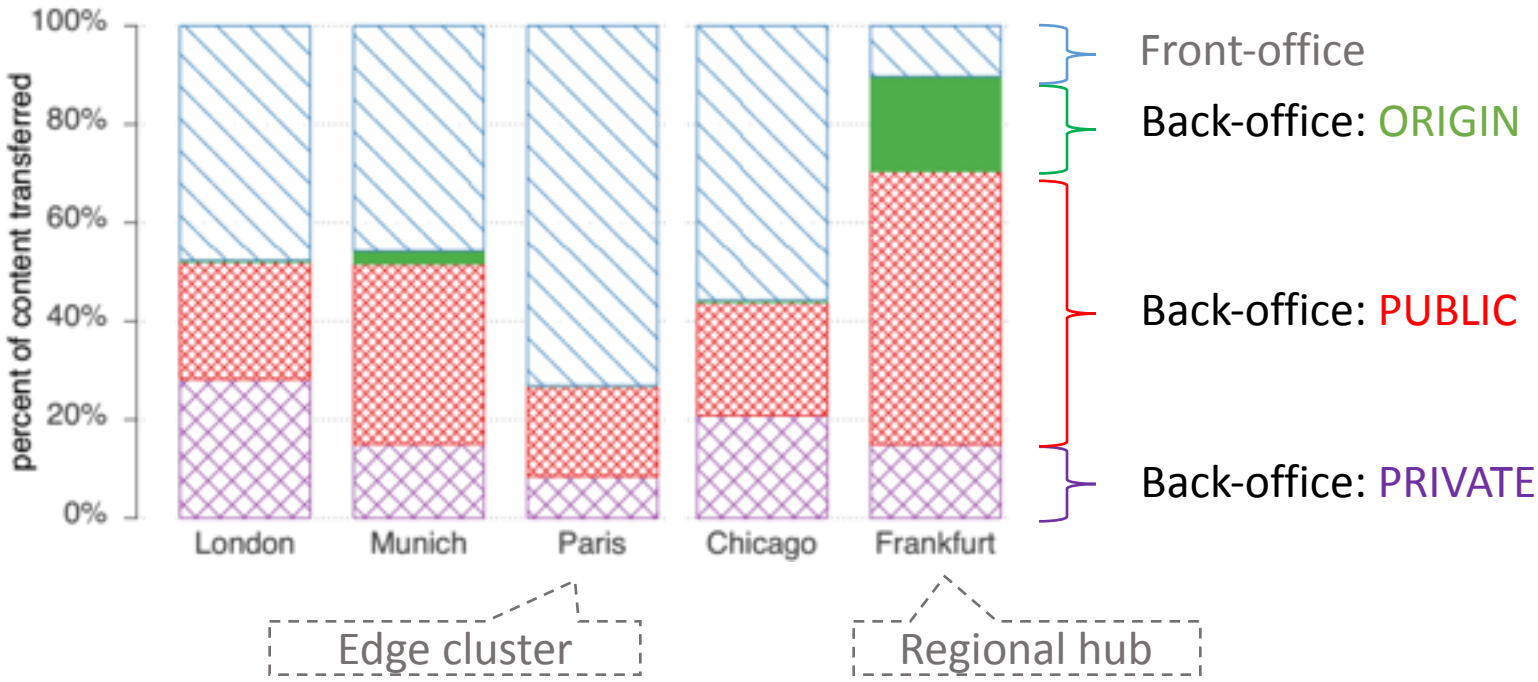

CDNs heavily rely on back-office traffic

Agenda

1. Introduction
2. Methodology and datasets
3. Characteristics
    1. Traffic
    2. Patterns
    3. Inter-domain perspective
4. CDN back-office traffic
5. The end-user perspective
6. Summary and implications

The end-user perspective

Residential broadband network: backbone latency (no access)

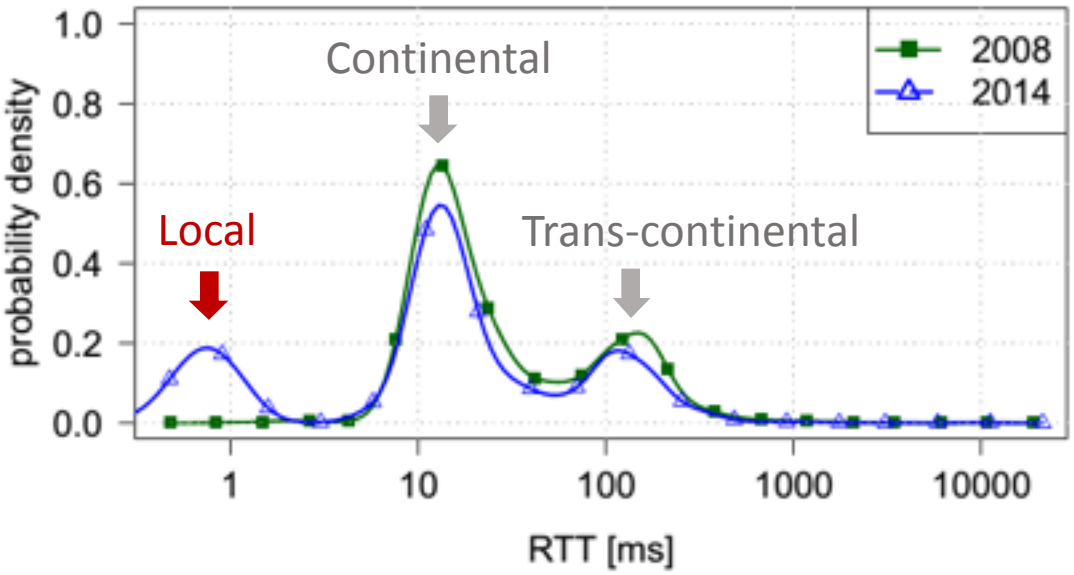

A smaller front-office: but the back-office may be large

Summary

1. A back-office to support the Web
2. Significant traffic: bytes and requests
3. Different type of traffic patterns
4. Visible at multiple peering links

An important yet understudied class of traffic

Implications

Feasibility to deploy new protocols:

- It is easier to change the back office than the front office

Implications

Feasibility to deploy new protocols:

- It is easier to change the back office than the front office

Performance evaluation:

- Interactions with the back office
- More users than anticipated

Implications

Feasibility to deploy new protocols:

- It is easier to change the back office than the front office

Performance evaluation:

- Interactions with the back office
- More users than anticipated

Opportunities:

- ISPs: micro-data centers, virtualized services
- IXPs: co-location strategies
- NSPs: new services e.g., SLAs

Back-office traffic on the Internet