

Fix with P6: Verifying Programmable Switches at Runtime



Apoorv Shukla
(Huawei Munich Research Center)



with K. N. Hudemann (SAP), Z. Vági (Swisscom), L. Hügerich (TU Berlin), G. Smaragdakis (TU Berlin/MPI), A. Hecker (Huawei MRC), S. Schmid (Vienna Uni.), and A. Feldmann (MPI)

Programming Protocol-Independent Packet Processors (P4²) Lang.

- High-level language: data plane programming
- P4 programs dictate: what packets and how packets processed
- Allows user-defined custom protocols/target independence
- P4 programs: compiled to run on P4 switches




[2] P. Bosshart, D. Daly, G. Gibby, M. Izzardy, N. McKeown, J. Rexford, C. Schlesinger, D. Talaycoy, A. Vahdat, G. Varghese, D. Walker. P4: Programming Protocol-Independent Packet Processors. SIGCOMM' 14.

Unfortunately, Runtime Issues happen

Runtime issues occur under diverse input workloads:

- Checksum and ECMP/hash-calculation
- Platform-dependent bugs



Need: Runtime verification to check switch/network behavior


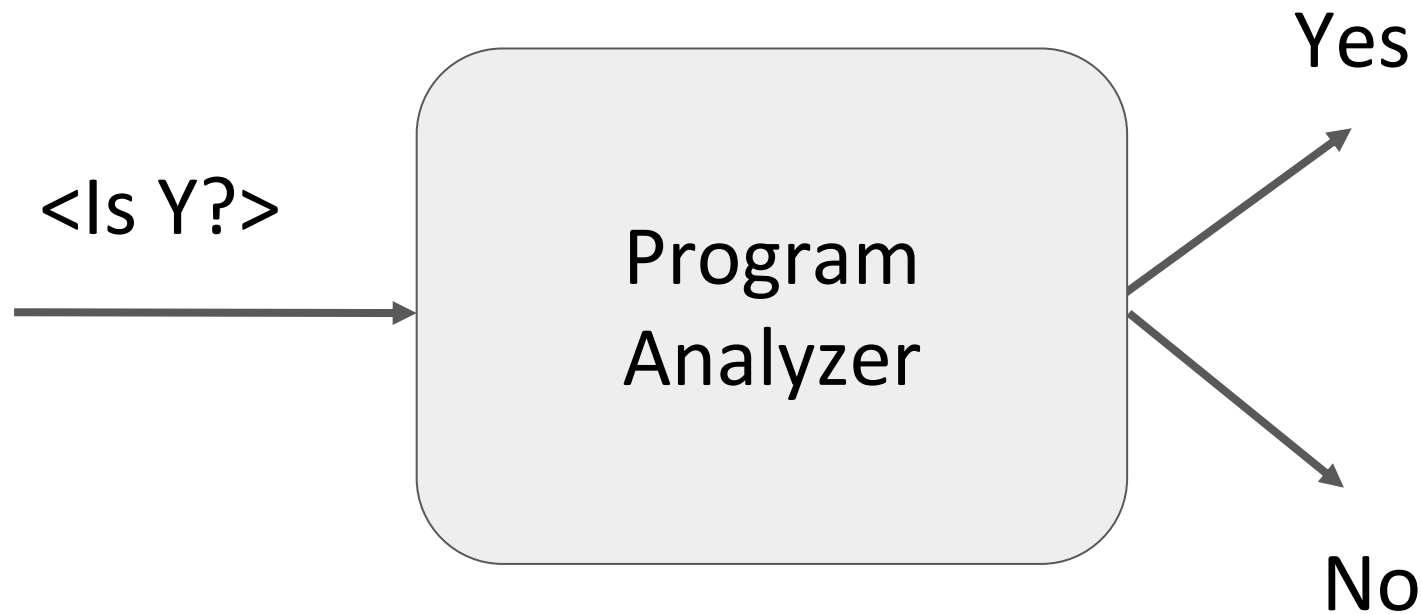
Runtime Bug Detection in Single P4 Switch: Hard

- P4 is *partial* program: rules populated at runtime
- Generate packets of interest: hard
- No runtime exceptions in P4: hard to catch
- Dense conditions, domain-specific constructs: hard to localize



Static analysis detects memory safety violations not runtime bugs

Challenge: Undecidability [Rice53]



Alan designed the perfect computer

Does Property Y occur or not?

Credit: <https://www.coopertoons.com/education/haltingproblem/haltingproblem.html>

Problem Statement

Can we detect “persistent” runtime bugs and mitigate them in a single P4 switch?

Progress in P4 Verification

P4NoD [MSR Tech. Report' 16]

P4K [P4 Workshop' 17]

Vera [SIGCOMM' 18]

p4v [SIGCOMM' 18]

ASSERT-P4 [SOSR'18, CoNEXT'18]

STATIC ANALYSIS

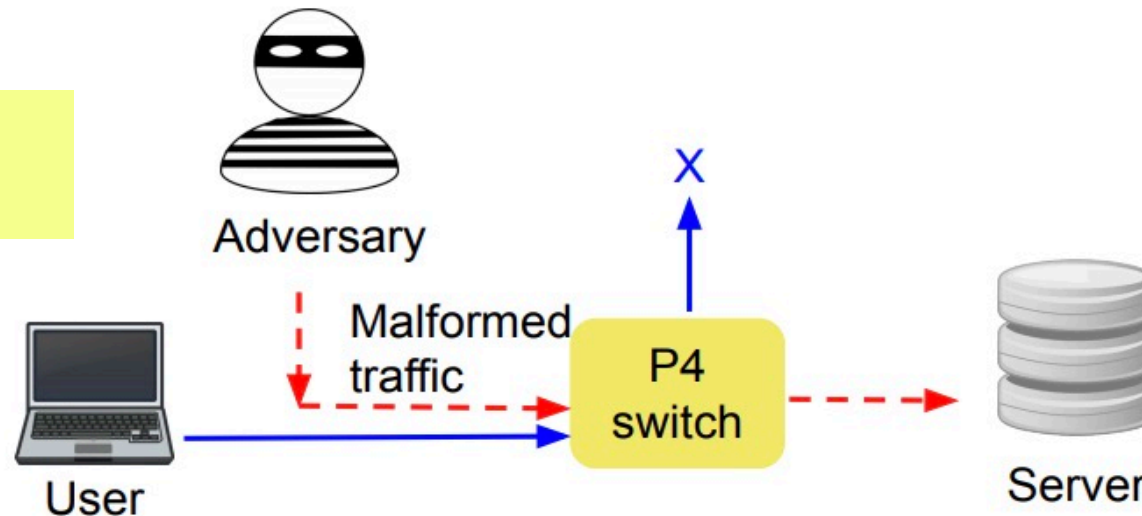
?

DYNAMIC ANALYSIS / RUNTIME VERIFICATION

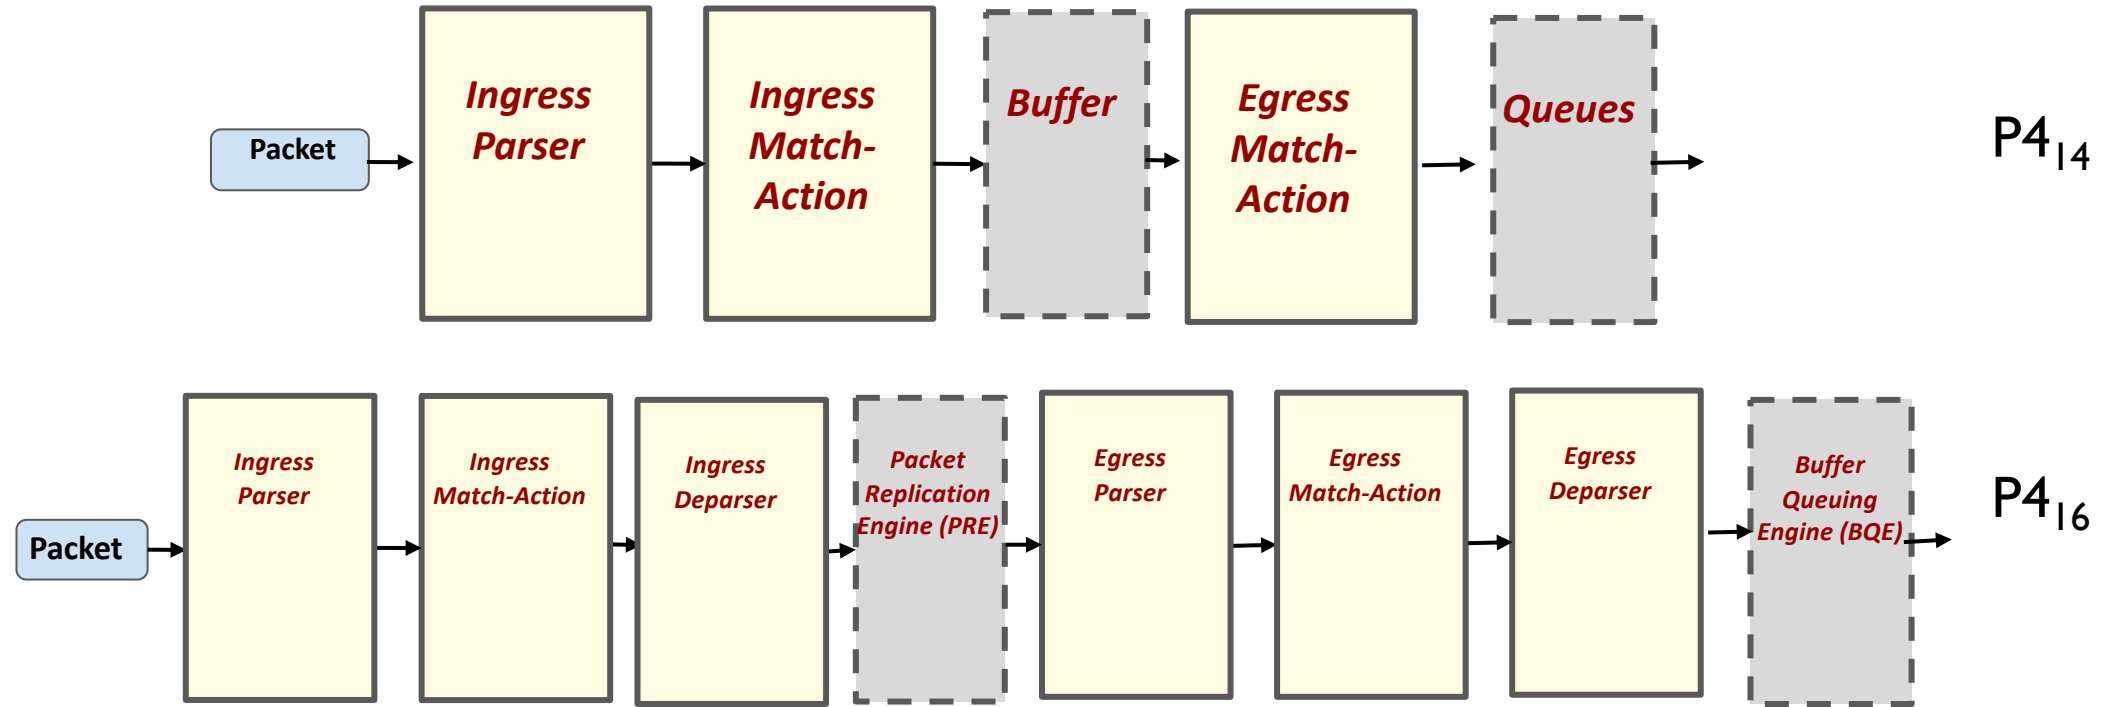
P4 Switch: Challenge- Platform-Independent Bug Example

- Packets with wrong IPv4 checksum is accepted (or generated)
- **Why:** Checksum calculation (or update) is not done on P4 switch

More Bug Examples in
Paper




Opportunities: Evolution of P4 Program Structure



Programmable blocks double from P4_{v4} to P4_{v6} are patchable!


P6: with Runtime Program Patching

- Get expected behavior from:
 1. p4q queries (test cases)
 2. Control plane configs (forwarding rules)
 3. Static analysis (reduce input space)
- (1) Generate packets: RL-guided fuzzing
- (2) Rewards if bug detected via query violation
- (3) Localize
- (4) Patch
- (5) Re-test



P6: Reinforcement Learning (RL) Agent for Guiding Fuzzing

- **States:** Byte-sequence in packet header
- **Actions:** Add/modify/delete bytes
- **Rewards:** $\begin{cases} 1, & \text{if packet triggered bug via query violation} \\ 0, & \text{otherwise} \end{cases}$



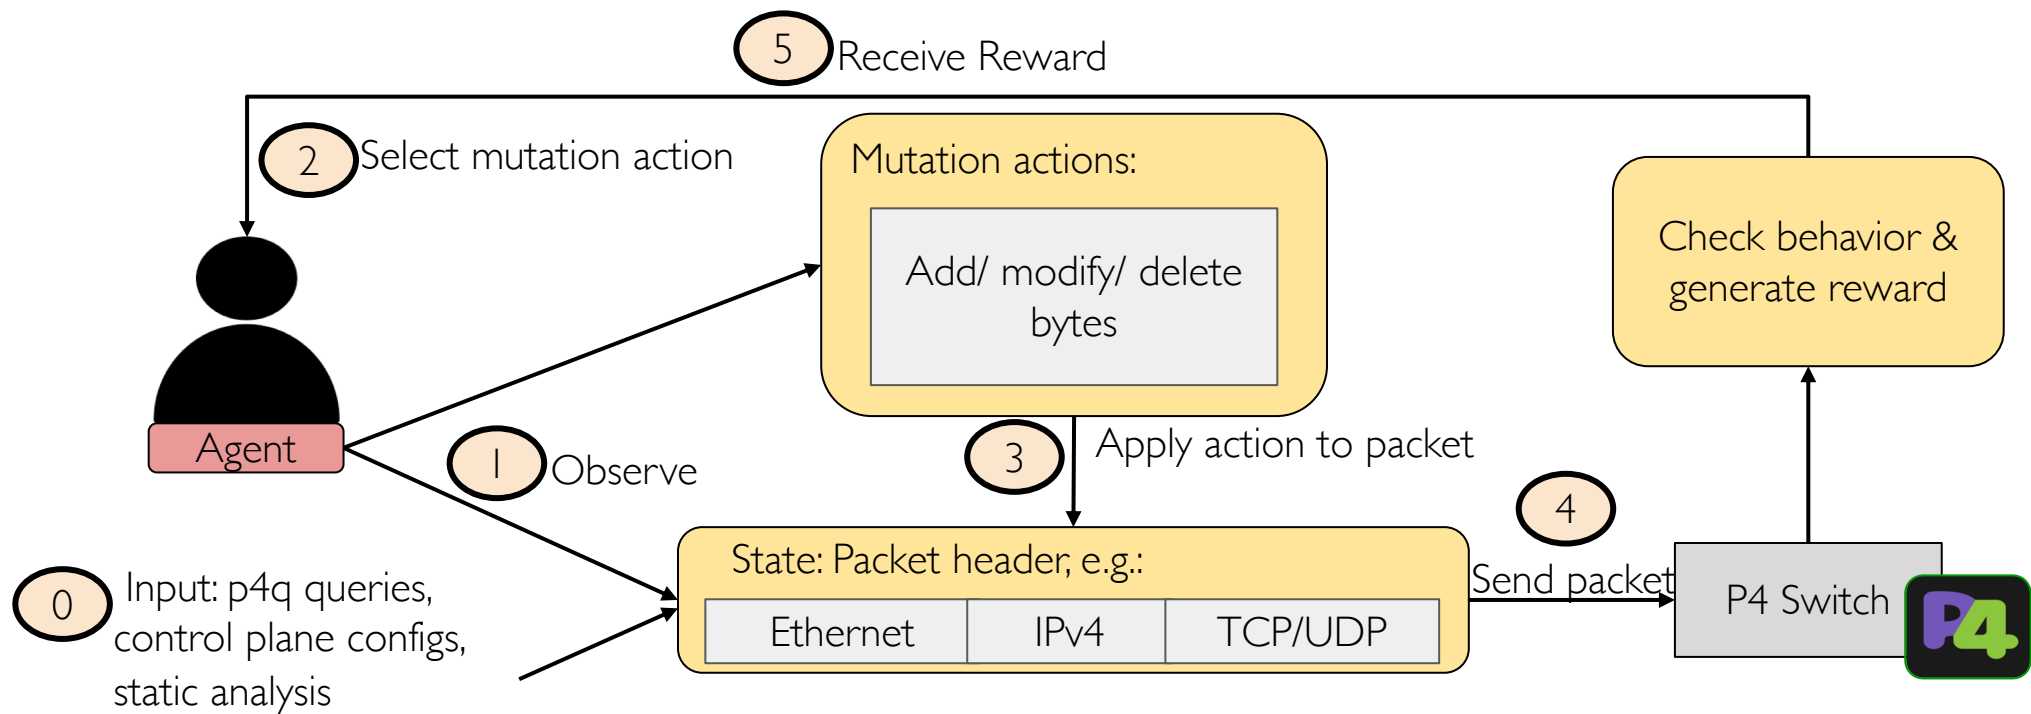
RL tries to maximise the cumulative reward (bugs)

Query Language: p4q

- Goal: Specify **expected** P4 switch behavior
- If- conditions to be fulfilled at ingress, then- conditions at egress, else- alternative conditions at egress
- Query violation signals bug

**(ing.hdr.ipv4 & ing.hdr.ipv4.version !=4,
egr.egress_port == False,)**

P6: Fuzzer (P6 Agent with RL-guided Fuzzing & Reward System)



RL tries to select appropriate action in a state to maximize cumulative reward

P6: P4Tarantula (Localize Detected Bugs in P4 Program)


- Tailor Tarantula for P4 programs
- Code traversal based on statements executed for an input
- Assign suspiciousness score (s) based on faulty statements
- Rank the statements as per s

```
parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start {
        transition parse_ethernet;
    }

    state parse_ethernet {
        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType) {
            TYPE_IPV4: parse_ipv4;
            default: accept;
        }
    }


    state parse_ipv4 {
        packet.extract(hdr.ipv4);
        transition accept;
    }
}
```



P6: Patcher (Patch, if Needed and Available)

Pre-generated library:

- Patches code lines without patches **and** high suspiciousness score violating p4q query
- Few lines, e.g., adding checks in parser
- Pass the sanity/regression tests



P6: In Action

P4 Source Code

```
.....  
state parse_ipv4 {  
    packet.extract(hdr.ipv4);  
    transition accept;  
}  
.....
```



P4 Source Code with Bugs

P4 Source Code

```
.....  
state parse_ipv4 {  
    packet.extract(hdr.ipv4);  
    transition accept;  
}  
.....
```



Localized by P4Tarantula

P4 Source Code

```
.....  
state parse_ipv4 {  
    packet.extract(hdr.ipv4);  
    verify(hdr.ipv4.version == 4, error.BadHeader);  
    verify(hdr.ipv4.ihl == 5, error.BadHeader);  
    verify(hdr.ipv4.len >= 20, error.BadHeader);  
    verify(hdr.ipv4.ttl >= 2, error.BadHeader);  
    transition accept; }  
.....  
apply {  
    if (standard_metadata.parser_error != error.NoError) {  
        mark_to_drop();  
        return;  
    }  
}  
.....
```



Patched by Patcher

P6 detects, localizes and patches bugs in a P4 program non-intrusively;
Important foray into self-driving networks

P6 performs Runtime Verification of P4 Switches

P4NoD [MSR Tech. Report' 16]
P4K [P4 Workshop' 17]
Vera [SIGCOMM' 18]
p4v [SIGCOMM' 18]
ASSERT-P4 [SOSR'18, CoNEXT'18]

STATIC ANALYSIS

P6 DYNAMIC ANALYSIS / RUNTIME VERIFICATION




P6: Evaluation Strategy

- 8 public P4 programs from P4.org, NetPaxos across 2 platforms
- Baselines:
 1. Advanced Agent: like P6 Agent, same mutation actions (Intelligent Baseline)
 2. IPv4-Based Agent: IPv4-based packets only
 3. Naïve Agent: Exhaustive generation

P6: Metrics


Gauge P6 performance over:

- # Bugs discovered
- P6 Performance: Detection Time
- P6 vs. Baselines:
 - Learning Performance: Rewards generated



P6 Evaluation: Bugs discovered quickly with few packets only

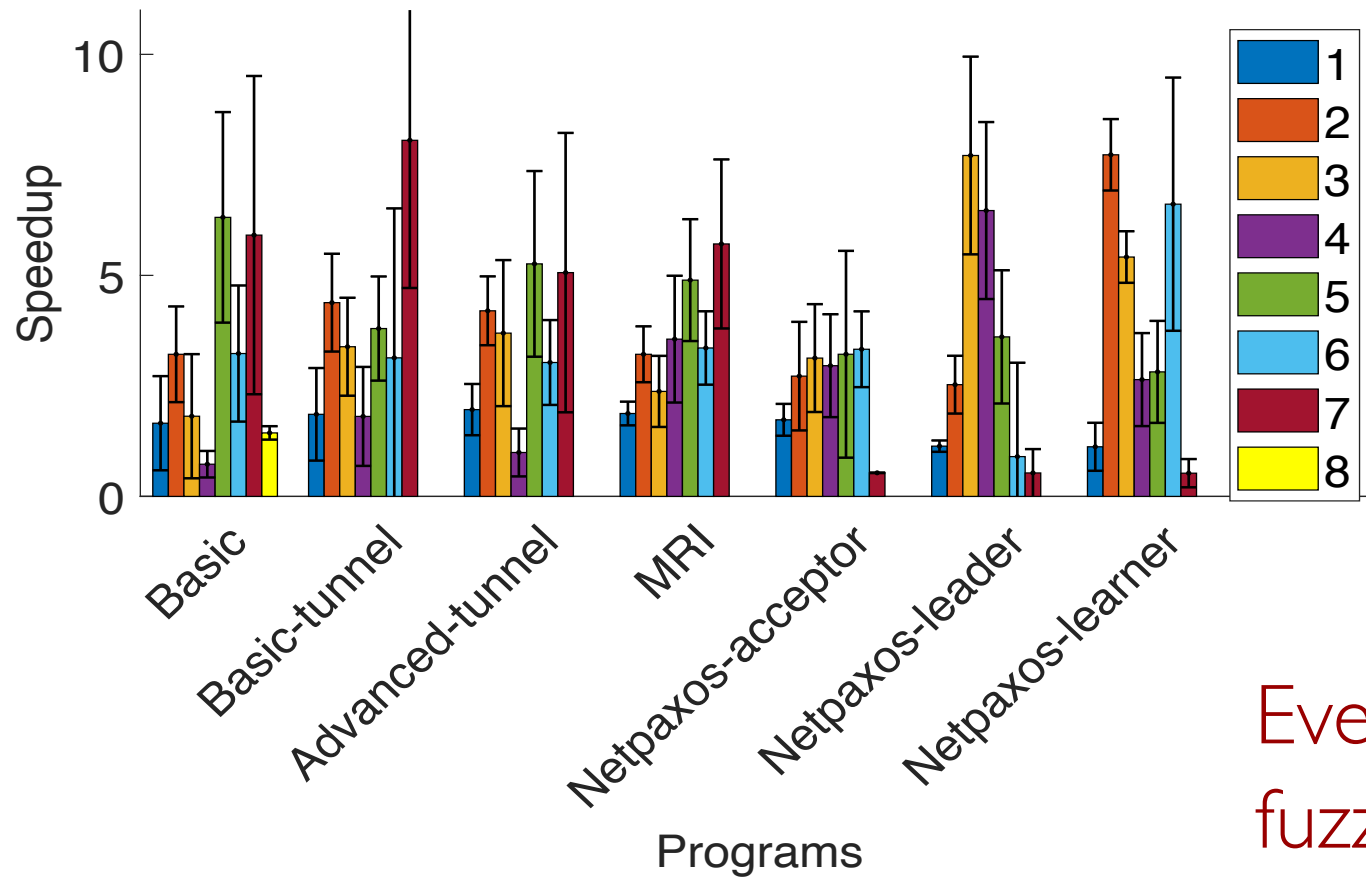
Bug IDs	Bugs
1	Accepted wrong checksum (PI)
2	Generated wrong checksum (PI)
3	Incorrect IP version (PI)
4	IP IHL value out of bounds (PI)
5	IP TotalLen value is too small (PI)
6	TTL 0 or 1 is accepted (PI)
7	TTL not decremented (PI)
8	Clone not dropped (PD)
9	Resubmitted packet not dropped (PD)
10	Multicast packet not dropped (PD)



PI – Platform-independent
PD – Platform-dependent

**P6 discovered 10 bugs in 8 publicly available P4 programs;
P6 detects bugs in ~2 secs in 7 programs,
in ~10 secs in switch.p4 (8715 LOC) with 28 packets only**

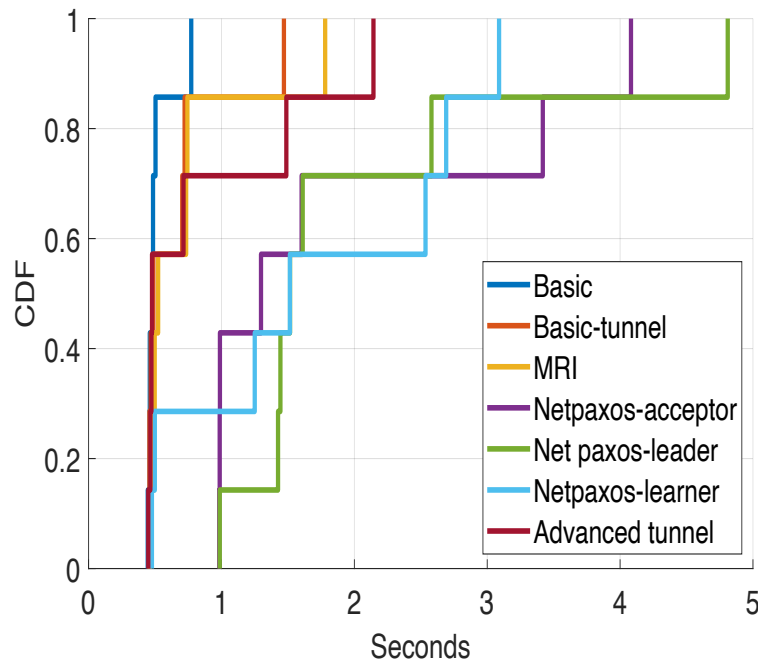
P6 Agent Detection Time Speedup vs. Advanced Agent



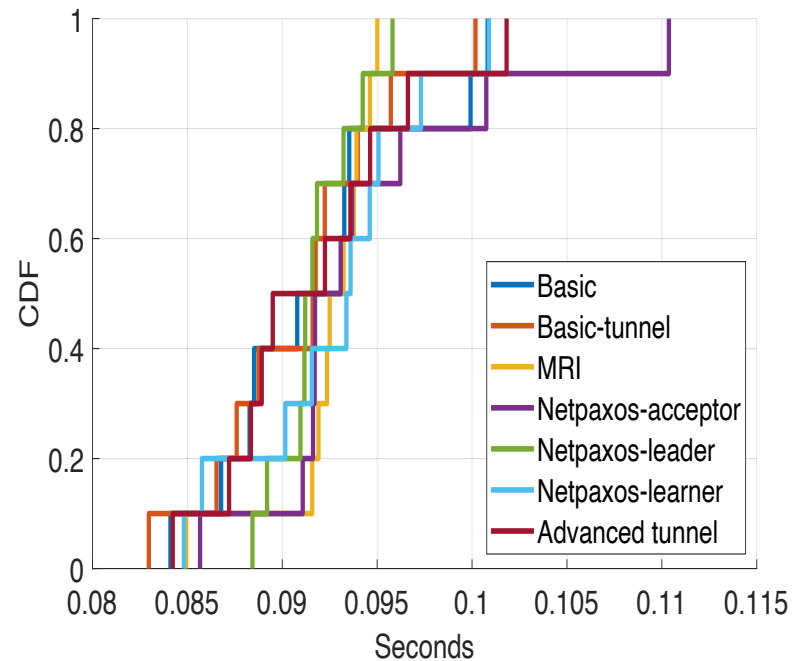
→ P6 up to 8× faster

Even faster on IPv4-based fuzzer, naïve cannot detect any


P6: Detection, Patching & Localization Times on Tofino Switch



Detection Time




Localization Time



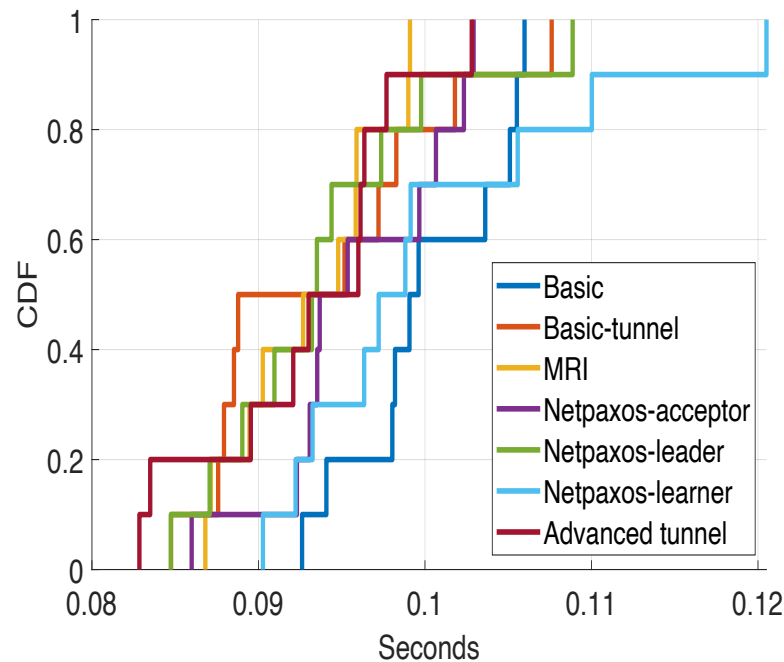
Patching Time

P6 detection time (second-scale), localization and patching time (ms-scale)
performance

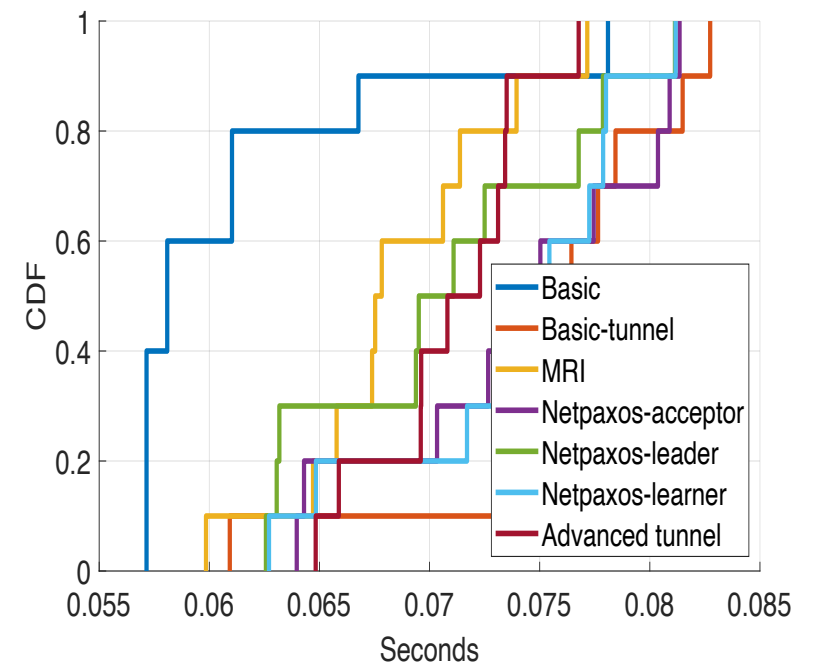
P6: Detection, Patching & Localization Times on bmv2 Switch



Detection Time



Localization Time



Patching Time

P6 detection time (second-scale), localization and patching time (ms-scale)


More Results in Paper

performance


Conclusion

- P6:
- enables **runtime verification** of P4 switches non-intrusively
 - uses static analysis- and machine learning-guided fuzzing to **detect** multiple runtime bugs. Later, **localizes** and **patches** with minimal human effort
 - **significantly outperforms baseline** bug detection approaches to detect platform-independent and -dependent bugs
 - foray into **Self-driving** or autonomous networks

THANK YOU!



P6



FUTURE
IS SELF-DRIVING

Contact: Dr. Apoorv Shukla
Email: apoorv.shukla@huawei.com